

Assignment 2 - Neural Networks

1.1 — Helper Functions

ReLU	<pre>def relu(S): X = np.copy(S) X[S<0] = 0 return X</pre>	S: a matrix of layer inputs - each row corresponds to a datapoint - each column corresponds to a neuron input at that layer
softmax	<pre>def softmax(S): X = np.exp(S) / np.sum(np.exp(S), axis=1, keepdims=True) return X</pre>	S: * see above
compute	<pre>def computeLayer(X, W, b): S = X @ W + b return S</pre>	X: a matrix of layer outputs W: weight matrix of a given layer b: bias vector of a given layer
averageCE	<pre>def avgCE(target, prediction): N = prediction.shape[0] L = - (1/N) * np.sum(target * np.log(prediction)) return L</pre>	target: a matrix of one-hot encoded labels prediction: a matrix of predicted labels (a probabilistic output)
gradCE	<pre>def gradCE(target, prediction): N = prediction.shape[0] dE_dX = - (1/N) * target / prediction return dE_dX</pre>	target: * see above prediction: * see above

For the gradCE function the average cross entropy loss was used.

Let x_{ij} denote the softmax of the outer layer predictions

Let y_{ij} denote the value of the one-hot encoded output at class j for datapoint i

$$\text{Then } E_{in} = -\frac{1}{N} \sum_n \sum_k y_{nk} \ln(x_{nk}) \implies \frac{\partial E_{in}}{\partial x_{ij}} = -\frac{1}{N} \frac{\partial}{\partial x_{ij}} y_{ij} \ln(x_{ij}) = -\frac{1}{N} \frac{y_{ij}}{x_{ij}}$$

$$\implies \frac{\partial E_{in}}{\partial X} = -\frac{1}{N} \frac{Y}{X}$$


* All these helper functions can be found in Appendix A

1.2 — Backpropagation Derivation

The following pages show how the vectorized forward and backpropagation algorithms are derived for this neural network

VECTORIZING THE NEURAL NETWORK

Model of L-layer NN



$$\text{Let } X^{(l)} = \begin{bmatrix} X_1^{(l)T} \\ \vdots \\ X_N^{(l)T} \end{bmatrix} \quad Y = \begin{bmatrix} y_1^T \\ \vdots \\ y_N^T \end{bmatrix}$$

$$\underline{b} = [b_1, \dots, b_k]$$

one-hot encoded

X_{ij} corresponds to the j^{th} feature of the i^{th} datapoint

$\dot{\underline{X}}$ indicates broadcasting.

FORWARD PROPAGATION VECTORIZATION

$$S_{nj}^{(l)} = \sum_i X_{ni}^{(l-1)} W_{ij}^{(l)} + b_j^{(l)}$$

$$X_{nj}^{(l)} = \theta(S_{nj}^{(l)})$$

$$\boxed{S^{(l)} = X^{(l-1)} W^{(l)} + \dot{\underline{b}}^{(l)}}$$

$$\boxed{X^{(l)} = \theta(S^{(l)})}$$

activation function of the l^{th} layer. (ie. ReLU or Softmax)

GRADIENT VECTORIZATION

Let $\delta_{nj}^{(l)} = \frac{\partial E_{in}}{\partial S_{nj}^{(l)}}$ the sensitivity of the i^{th} datapoint at the node j layer l .

$$\frac{\partial E_{in}}{\partial W_{ij}^{(l)}} = \sum_n \frac{\partial E_{in}}{\partial S_{nj}^{(l)}} \cdot \frac{\partial S_{nj}^{(l)}}{\partial W_{ij}^{(l)}}$$

(This is b/c $E_{in} = \frac{1}{N} \sum_n E_n = f(S_1^{(l)}, S_2^{(l)}, \dots, S_N^{(l)})$ which means we need to apply a branching chain rule)

$$= \sum_n \delta_{nj}^{(l)} \cdot X_{ni}^{(l-1)} = \sum_n (X_{in}^{(l-1)})^T \cdot \delta_{nj}^{(l)} \Rightarrow$$

$$\boxed{\frac{\partial E_{in}}{\partial W^{(l)}} = [X^{(l-1)}]^T \delta^{(l)}}$$

$$\frac{\partial E_{in}}{\partial b_j^{(l)}} = \sum_n \frac{\partial E_{in}}{\partial S_{nj}^{(l)}} \cdot \frac{\partial S_{nj}^{(l)}}{\partial b_j^{(l)}}$$

$$= \sum_n \delta_{nj}^{(l)} \cdot 1$$

$$\Rightarrow \boxed{\frac{\partial E_{in}}{\partial \underline{b}^{(l)}} = \underline{1}^T \delta^{(l)}}$$

BACKPROPAGATION VECTORIZATION

$$\begin{aligned}\delta^{(l)} &= \frac{\partial E_{in}}{\partial S_{nj}^{(l)}} = \frac{\partial E_{in}}{\partial X_{nj}^{(l)}} \cdot \frac{\partial X_{nj}^{(l)}}{\partial S_{nj}^{(l)}} \\ &= \left(\sum_k \frac{\partial E_{in}}{\partial S_{nk}^{(l+1)}} \cdot \frac{\partial S_{nk}^{(l+1)}}{\partial X_{nj}^{(l)}} \right) \cdot \theta'(S_{nj}^{(l)}) \\ &= \left(\sum_k \delta_{nk}^{(l+1)} w_{jk}^{(l+1)} \right) \cdot \theta'(S_{nj}^{(l)}) = \left(\sum_k \delta_{nk}^{(l+1)} (w_{kj}^{(l+1)})^T \right) \cdot \theta'(S_{nj}^{(l)})\end{aligned}$$

$$\Rightarrow \boxed{\delta^{(l)} = \left(\delta^{(l+1)} [W^{(l+1)}]^T \right) \otimes \theta'(S^{(l)})}$$


in the hidden layer for Relu.
 $\theta'(S^{(l)}) = \text{sign}(S^{(l)})$

SEEDING the SENSITIVITY

What is $\delta^{(L)}$? $L = \text{output layer}$.

$$\delta_{nj}^{(L)} = \frac{\partial E_{in}}{\partial S_{nj}^{(L)}} = \sum_k \frac{\partial E_{in}}{\partial X_{nk}^{(L)}} \cdot \frac{\partial X_{nk}^{(L)}}{\partial S_{nj}^{(L)}}$$

(Note: unlike for hidden layer activation, the softmax depends on all $S_{nj}^{(L)}$ $\rightarrow \therefore$ we need branching chain rule)



* Notice that the softmax activation depends on all $S_j^{(L)}$

$$E_{in} = -\frac{1}{N} \sum_{n=1}^N \sum_{k=1}^K y_{nk} \ln(X_{nk}^{(L)}) \quad \text{for part 1.1.5}$$

$$\frac{\partial E_{in}}{\partial X_{nk}^{(L)}} = -\frac{1}{N} \frac{y_{nk}}{X_{nk}^{(L)}} \Rightarrow \frac{\partial E_{in}}{\partial X^{(L)}} = -\frac{1}{N} \frac{Y}{X^{(L)}}$$

$$\frac{\partial X_{nk}^{(L)}}{\partial S_{nj}^{(L)}} = \frac{\partial}{\partial S_{nj}^{(L)}} \left(\sigma(S_{nk}^{(L)}) \right) = X_{nk}^{(L)} (\delta_{kj} - X_{nj}^{(L)})$$

\uparrow dirac delta.

$$\delta_{nj}^{(L)} = \sum_k \frac{\partial E_{in}}{\partial X_{nk}^{(L)}} \cdot \frac{\partial X_{nk}^{(L)}}{\partial S_{nj}^{(L)}}$$

$$= -\frac{1}{N} \sum_k \frac{y_{nk}}{X_{nk}^{(L)}} \cdot X_{nk}^{(L)} (\delta_{kj} - X_{nj}^{(L)})$$

$$= -\frac{1}{N} \sum_k (y_{nk} \delta_{kj} - y_{nk} X_{nj}^{(L)})$$

$$\Rightarrow = -\frac{1}{N} \left(\sum_k y_{nk} \delta_{kj} - X_{nj}^{(L)} \sum_k y_{nk} \right)$$

$$\Rightarrow \delta^{(L)} = -\frac{1}{N} (Y \cdot I - X^{(L)})$$

is always = 1 b/c one-hot encoded.

$$\boxed{\delta^{(L)} = \frac{1}{N} (X^{(L)} - Y)}$$

For the specific 2-Layer Neural Network we are asked to implement...



$$\left\{ \begin{array}{l} \frac{\partial L}{\partial w_0} = \frac{\partial E_{in}}{\partial w^{(2)}} = [X^{(1)}]^T \delta^{(2)} \\ \frac{\partial L}{\partial b_0} = \frac{\partial E_{in}}{\partial b^{(2)}} = \underline{1}^T \delta^{(2)} \end{array} \right.$$

$$\left\{ \begin{array}{l} \frac{\partial L}{\partial w_n} = \frac{\partial E_{in}}{\partial w^{(1)}} = [X^{(0)}]^T \delta^{(1)} \\ \frac{\partial L}{\partial b_n} = \frac{\partial E_{in}}{\partial b^{(1)}} = \underline{1}^T \delta^{(1)} \end{array} \right.$$

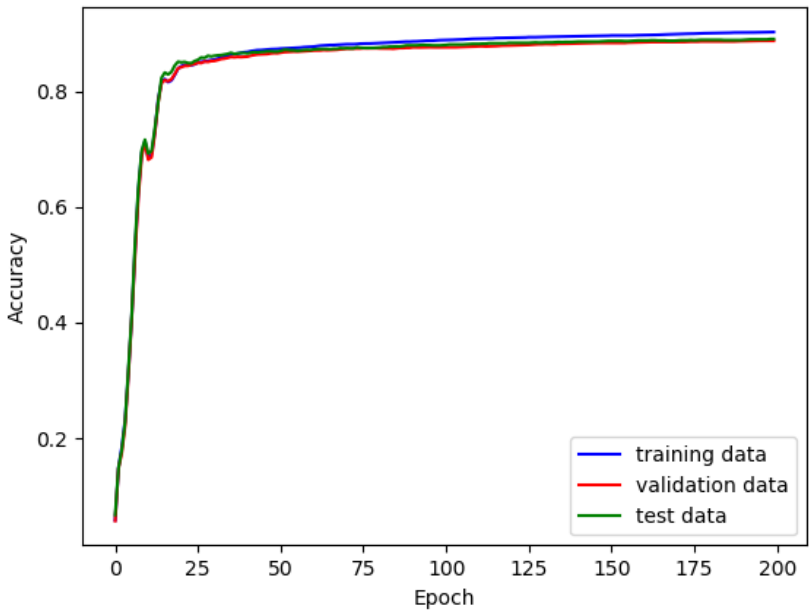
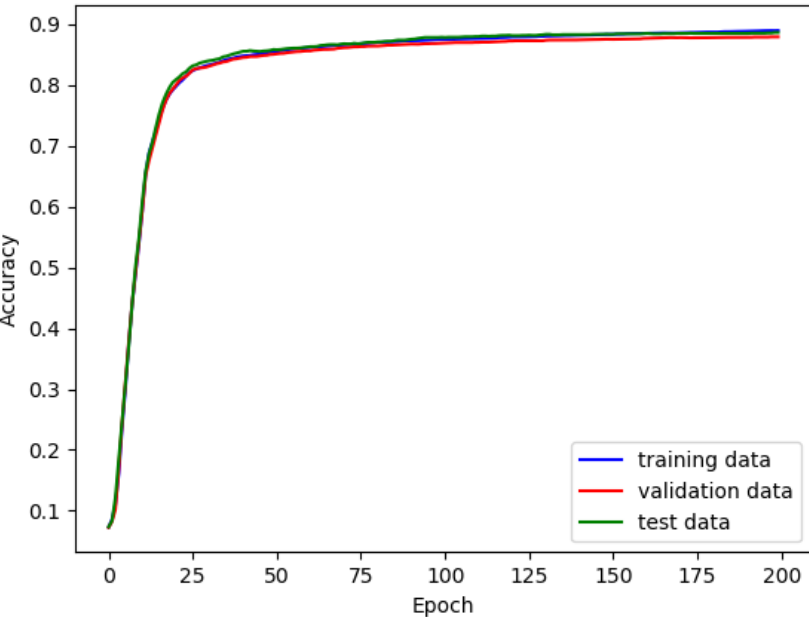
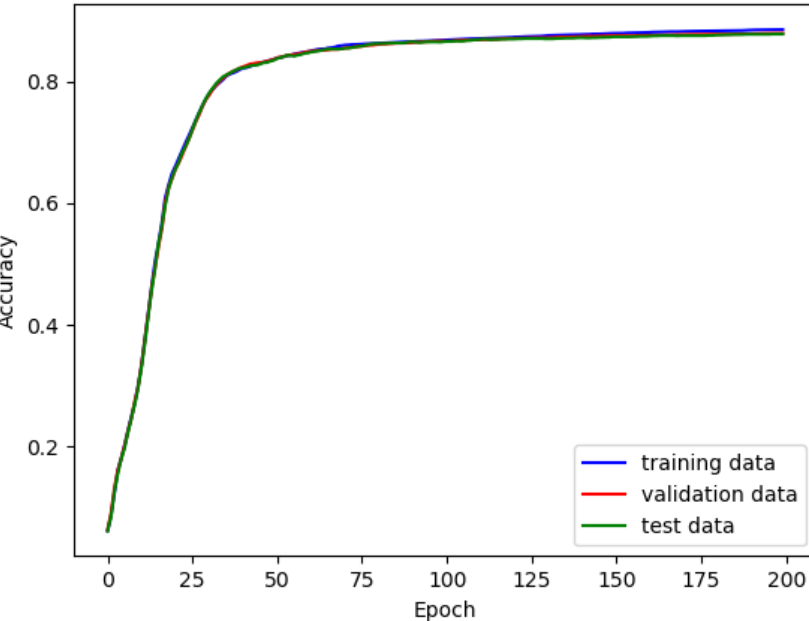
$$\left\{ \begin{array}{l} \delta^{(2)} = \frac{1}{N} (X^{(2)} - Y) \\ \delta^{(1)} = (\delta^{(2)} [W^{(2)}]^T) \otimes \text{Sign}(s^{(1)}) \end{array} \right.$$

1.3 — Learning

- * all code for learning of the neural network can be found in Appendix B — `gradient_descent()`
- * Appendix B — `main()` executes the NN learning with learning rate = 0.005, epochs = 200, hidden units = 1000



1.4.1 — Number of Hidden Neurons


2000 Neurons	<p>Accuracy Curves</p>  <p>The graph shows accuracy on the y-axis (0.2 to 0.8) and epoch on the x-axis (0 to 200). Three lines represent training data (blue), validation data (red), and test data (green). All three lines rise sharply from epoch 0 to 25, reaching an accuracy of approximately 0.85. After epoch 25, the lines plateau, with training data reaching about 0.90 and validation/test data reaching about 0.88 by epoch 200.</p>	<p>TRAINING</p> <p>Loss: 0.3697683805018834 Accuracy: 0.9023</p> <p>VALIDATION</p> <p>Loss: 0.41339738342300447 Accuracy: 0.8871666666666667</p> <p>TESTING</p> <p>Loss: 0.4193425232371352 Accuracy: 0.8902349486049926</p>
500 Neurons	<p>Accuracy Curves</p>  <p>The graph shows accuracy on the y-axis (0.1 to 0.9) and epoch on the x-axis (0 to 200). Three lines represent training data (blue), validation data (red), and test data (green). All three lines rise sharply from epoch 0 to 25, reaching an accuracy of approximately 0.85. After epoch 25, the lines plateau, with training data reaching about 0.89 and validation/test data reaching about 0.88 by epoch 200.</p>	<p>TRAINING</p> <p>Loss: 0.4139265238179206 Accuracy: 0.8895</p> <p>VALIDATION</p> <p>Loss: 0.44251219875802256 Accuracy: 0.879</p> <p>TESTING</p> <p>Loss: 0.443531155268261 Accuracy: 0.8861967694566814</p>
100 Neurons	<p>Accuracy Curves</p>  <p>The graph shows accuracy on the y-axis (0.2 to 0.8) and epoch on the x-axis (0 to 200). Three lines represent training data (blue), validation data (red), and test data (green). All three lines rise sharply from epoch 0 to 25, reaching an accuracy of approximately 0.85. After epoch 25, the lines plateau, with training data reaching about 0.89 and validation/test data reaching about 0.88 by epoch 200.</p>	<p>TRAINING</p> <p>Loss: 0.4322297150343103 Accuracy: 0.8851</p> <p>VALIDATION</p> <p>Loss: 0.4584138791927889 Accuracy: 0.8791666666666667</p> <p>TESTING</p> <p>Loss: 0.46465170276411927 Accuracy: 0.8777533039647577</p>

Analysis

By comparing the training accuracies of neural networks with different number of hidden units, we noticed that the more number of hidden units, the lower the total loss and higher the classification accuracy. This is also true when we compared the test losses and accuracies, although in other cases this may not be true due to overfitting.

1.4.2 — Early Stopping

The early stopping point is the point which has the best test accuracy. In other words, the maximum of the test accuracy curve. Examining the zoomed in accuracy plot from 1.3 we can approximate the early stopping epoch.



The vertical blue line indicates the approximate maximum test accuracy. This occurs at **epoch = 140**

The approximate classification accuracies at epoch = 140 are:

Training Accuracy: 0.883
Validation Accuracy 0.875
Test Accuracy: 0.880



2.1 — Tensorflow CNN Model Implementation

* all code for convolutional neural network model can be found in Appendix C — `build_CNN()`

2.2 — CNN Training

* all code for learning of the CNN can be found in Appendix C — `train_CNN()`

* Appendix C — `main()` executes the NN learning with given parameters (set `weight_decay=0` and `p=1` for no regularization)



2.3.1 — L2 Regularization

- * all code for CNN regularization can be found in Appendix C — main()
- * set weight_decay = 0.01 or 0.1 or 0.5 and p=1

weight decay	Final Classification Accuracy		
	Training Data	Validation Data	Test Data
0.01	0.9952	0.93133336	0.9295154
0.1	0.946	0.9206667	0.92217326
0.5	0.8908	0.8843333	0.89133626


The L2 Regularization parameter directly impacts the spread of the classification accuracy between training, validation, and test datasets.

- At low weight decay (0.01) the model is overfitted which causes the training accuracy to be much higher than the validation and test accuracies
- At high weight decay (0.5) the model may be under-fitted, but the difference between the training and test accuracies are much smaller than with a low weight decay

For the given dataset, the test accuracy actually decreases with weight decay greater than 0.1. This suggests that the model may be over-regularized. A more optimal decay value may be found between 0.01 and 0.1.


2.3.2 — Dropout

- * all code for CNN dropout can be found in Appendix C — main()
- * set weight_decay = 0 and p = 0.9 or 0.75 or 0.5



p=0.75

Accuracy Curves



TRAINING

Loss: 0.002391771
Accuracy: 0.9994

VALIDATION


Loss: 0.46718428
Accuracy: 0.9316667

TESTING

Loss: 0.49651933
Accuracy: 0.9298825

p=0.5

Accuracy Curves



TRAINING

Loss: 0.004340526
Accuracy: 0.9992


VALIDATION

Loss: 0.45040885
Accuracy: 0.93266666

TESTING

Loss: 0.48736116
Accuracy: 0.9265786

p=0.1



TRAINING

Loss: 0.015503812
Accuracy: 0.9955

VALIDATION

Loss: 0.27441043
Accuracy: 0.939

TESTING

Loss: 0.27909115
Accuracy: 0.9390602

* Dropout does not seem very effective when applied to only one fully connected layer in this CNN

Appendix A — starter.py

```
# Implementation of a neural network using only Numpy
# - trained using gradient descent with momentum
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import time
import os
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'

# Load the data
def loadData():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx] / 255.0
        Target = Target[randIndx]
        trainData, trainTarget = Data[:10000], Target[:10000]
        validData, validTarget = Data[10000:16000], Target[10000:16000]
        testData, testTarget = Data[16000:], Target[16000:]
    return trainData, validData, testData, trainTarget, validTarget, testTarget

def parseData(data):
    num_data = data.shape[0]
    X = data.reshape(num_data, -1)
    return X

def convertOneHot(trainTarget, validTarget, testTarget):
    newtrain = np.zeros((trainTarget.shape[0], 10))
    newvalid = np.zeros((validTarget.shape[0], 10))
    newtest = np.zeros((testTarget.shape[0], 10))

    for item in range(0, trainTarget.shape[0]):
        newtrain[item][trainTarget[item]] = 1
    for item in range(0, validTarget.shape[0]):
        newvalid[item][validTarget[item]] = 1
    for item in range(0, testTarget.shape[0]):
        newtest[item][testTarget[item]] = 1
    return newtrain, newvalid, newtest

# PARSE THE DATA -----
trainData, validData, testData, trainTarget, validTarget, testTarget = loadData()
X_train, X_valid, X_test = parseData(trainData), parseData(validData), parseData(testData)
Y_train, Y_valid, Y_test = convertOneHot(trainTarget, validTarget, testTarget)

def shuffle(trainData, trainTarget):
    np.random.seed(421)
    randIndx = np.arange(len(trainData))
    target = trainTarget
    np.random.shuffle(randIndx)
    data, target = trainData[randIndx], target[randIndx]
    return data, target

def relu(S):
    X = np.copy(S)
    X[S<0] = 0
    return X
```

```

def derivative_relu(S):
    dS = np.zeros_like(S)
    dS[S>0] = 1
    return dS

def softmax(S):
    X = np.exp(S) / np.sum(np.exp(S), axis=1, keepdims=True)
    return X

def computeLayer(X, W, b):
    S = X @ W + b
    return S

def avgCE(target, prediction):
    N = prediction.shape[0]
    L = - (1/N) * np.sum(target * np.log(prediction))
    return L

def gradCE(target, prediction):
    N = prediction.shape[0]
    dE_dX = - (1/N) * target / prediction
    return dE_dX

def accuracy(Y, Y_pred):
    j = np.argmax(Y_pred, axis=1)
    i = np.arange(Y.shape[0])
    return np.mean(Y[i, j])

```

Appendix B — neural_network_numpy.py

```
from starter import *
import time

def xavier_init(neurons_in, n_units, neurons_out):
    shape = (neurons_in, n_units)
    var = 2./(neurons_in + neurons_out)
    W = np.random.normal(0, np.sqrt(var), shape)
    return W

def init_weights(n_input, n_hidden, n_output):
    W = []
    W.append(None)
    W.append(xavier_init(n_input, n_hidden, n_output))
    W.append(xavier_init(n_hidden, n_output, 1))
    return W

def init_biases(n_input, n_hidden, n_output):
    b = []
    b.append(None)
    b.append(np.zeros((1, n_hidden)))
    b.append(np.zeros((1, n_output)))
    return b

def forward_propagation(X_input, W, b):
    X, S = [None]*3, [None]*3
    X[0] = X_input

    # UPDATE HIDDEN LAYER
    S[1] = X[0] @ W[1] + b[1]
    X[1] = relu(S[1])

    # UPDATE OUTPUT LAYER
    S[2] = X[1] @ W[2] + b[2]
    X[2] = softmax(S[2])

    return X, S

def backpropagation(X, S, W, Y):
    SENS = [None]*3

    # SEED SENSITIVITY
    N = Y.shape[0]
    SENS[2] = (1/N) * (X[2] - Y)

    # BACKPROPAGATION
    SENS[1] = (SENS[2] @ (W[2]).T) * derivative_relu(S[1])

    return SENS

def compute_gradients(X_input, Y, W, b):
    gradW = [0]*3
    gradb = [0]*3
    N = X_input.shape[0]

    # RUN PROPAGATIONS
    X, S = forward_propagation(X_input, W, b)
    SENS = backpropagation(X, S, W, Y)
```

```
# GRADIENT of OUTPUT LAYER WEIGHTS + BIASES
```

```
gradW[2] = (X[1]).T @ SENS[2]
```

```
gradb[2] = np.sum(SENS[2], axis=0)
```

```
# GRADIENT of HIDDEN LAYER WEIGHTS + BIASES
```

```
gradW[1] = (X[0]).T @ SENS[1]
```

```
gradb[1] = np.sum(SENS[1], axis=0)
```

```
return gradW, gradb
```

```
def measure_performance(W, b):
```

```
Y_pred, S = forward_propagation(X_train, W, b)
```

```
train_loss = avgCE(Y_train, Y_pred[2])
```

```
train_acc = accuracy(Y_train, Y_pred[2])
```

```
Y_pred, S = forward_propagation(X_valid, W, b)
```

```
valid_loss = avgCE(Y_valid, Y_pred[2])
```

```
valid_acc = accuracy(Y_valid, Y_pred[2])
```

```
Y_pred, S = forward_propagation(X_test, W, b)
```

```
test_loss = avgCE(Y_test, Y_pred[2])
```

```
test_acc = accuracy(Y_test, Y_pred[2])
```

```
return train_loss, train_acc, valid_loss, valid_acc, test_loss, test_acc
```

```
def gradient_descent(X, Y, n_epochs, alpha, gamma, n_hidden_units):
```

```
n_input_neurons = X.shape[1]
```

```
n_output_neurons = Y.shape[1]
```

```
# INITIALIZE WEIGHTS + BIASES
```

```
W = init_weights(n_input_neurons, n_hidden_units, n_output_neurons)
```

```
b = init_biases(n_input_neurons, n_hidden_units, n_output_neurons)
```

```
VW_o = np.ones_like(W[2]) * 1e-5
```

```
VW_h = np.ones_like(W[1]) * 1e-5
```

```
Vb_o = np.ones_like(b[2]) * 1e-5
```

```
Vb_h = np.ones_like(b[1]) * 1e-5
```

```
# Create Loss/Accuracy dictionaries
```

```
loss = {'train': [], 'valid': [], 'test': []}
```

```
accuracy = {'train': [], 'valid': [], 'test': []}
```

```
for t in range(n_epochs):
```

```
    gradW, gradb = compute_gradients(X, Y, W, b)
```

```
    print("EPOCH {}".format(t))
```

```
    # UPDATE OUTPUT LAYER
```

```
    VW_o = gamma * VW_o + alpha * gradW[2]
```

```
    W[2] = W[2] - VW_o
```

```
    Vb_o = gamma * Vb_o + alpha * gradb[2]
```

```
    b[2] = b[2] - Vb_o
```

```
    # UPDATE HIDDEN LAYER
```

```
    VW_h = gamma * VW_h + alpha * gradW[1]
```

```
    W[1] = W[1] - VW_h
```

```
    Vb_h = gamma * Vb_h + alpha * gradb[1]
```

```
    b[1] = b[1] - Vb_h
```

```
    # MEASURE PERFORMANCE
```

```
    train_loss, train_acc, valid_loss, valid_acc, test_loss, test_acc = measure_performance(W, b)
```

```
    loss['train'].append(train_loss)
```

```
    accuracy['train'].append(train_acc)
```

```
    loss['valid'].append(valid_loss)
```

```
    accuracy['valid'].append(valid_acc)
```

```

    loss['test'].append(test_loss)
    accuracy['test'].append(test_acc)

return W, b, loss, accuracy

def main():
    start_time = time.time()
    W, b, loss, accuracy = gradient_descent(X_train, Y_train,
                                             n_epochs=200,
                                             alpha=0.005,
                                             gamma=0.9,
                                             n_hidden_units=100)
    end_time = time.time()
    print("--- %s seconds ---" % (time.time() - start_time))

    print("TRAINING -----")
    print("Loss: ", loss['train'][-1])
    print("Accuracy:", accuracy['train'][-1])

    print("VALIDATION -----")
    print("Loss: ", loss['valid'][-1])
    print("Accuracy:", accuracy['valid'][-1])

    print("TESTING -----")
    print("Loss: ", loss['test'][-1])
    print("Accuracy:", accuracy['test'][-1])

    plt.plot(loss['train'], color='blue', label='training data')
    plt.plot(loss['valid'], color='red', label='validation data')
    plt.plot(loss['test'], color='green', label='test data')
    plt.legend()
    plt.title('Loss Curves')
    plt.ylabel('Average CE Loss')
    plt.xlabel('Epoch')
    plt.show()

    plt.plot(accuracy['train'], color='blue', label='training data')
    plt.plot(accuracy['valid'], color='red', label='validation data')
    plt.plot(accuracy['test'], color='green', label='test data')
    plt.legend()
    plt.title('Accuracy Curves')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.show()

if __name__ == '__main__':
    main()

```


Appendix C — neural_network_tensorflow.py

```
from starter import *
import time

img_rows = 28
img_cols = 28
img_depth = 1
n_features = 784
n_classes = 10

def convolution_layer(input, n_channels, filter_size, n_filters, name):
    weights = tf.get_variable(
        "weight_{}".format(name),
        shape=[filter_size, filter_size, n_channels, n_filters],
        initializer=tf.contrib.layers.xavier_initializer()
    )
    biases = tf.get_variable(
        "bias_{}".format(name),
        shape=[n_filters],
        initializer=tf.constant_initializer(0.0)
    )
    layer = tf.nn.conv2d(
        input=input,
        filter=weights,
        strides=[1, 1, 1, 1],
        padding="SAME"
    )

    layer = layer + biases
    return layer, weights, biases

def relu_layer(input):
    layer = tf.nn.relu(input)
    return layer

def batch_normalization_layer(input):
    batch_mean, batch_var = tf.nn.moments(input, axes=[0, 1, 2])
    scale = tf.Variable(tf.ones([32]))
    beta = tf.Variable(tf.zeros([32]))

    layer = tf.nn.batch_normalization(input, batch_mean, batch_var,
        offset=beta,
        scale=scale,
        variance_epsilon=1e-3
    )
    return layer

def pooling_layer(input, kernel_size, stride_size):
    layer = tf.nn.max_pool(
        value=input,
        ksize=[1, kernel_size, kernel_size, 1],
        strides=[1, stride_size, stride_size, 1],
        padding="SAME"
    )
    return layer

def flatten_layer(input):
    layer = tf.contrib.layers.flatten(input)
    return layer
```

```

def fully_connected_layer(input, n_inputs, n_outputs, name):
    weights = tf.get_variable(
        "weight_{}".format(name),
        shape=[n_inputs, n_outputs],
        initializer=tf.contrib.layers.xavier_initializer()
    )
    biases = tf.get_variable(
        "bias_{}".format(name),
        shape=[n_outputs],
        initializer=tf.constant_initializer(0.0)
    )

    layer = tf.matmul(input, weights) + biases
    return layer, weights, biases

def build_CNN(learning_rate, weight_decay):
    # Define placeholders for input
    X = tf.placeholder(tf.float32, shape=[None, n_features], name="X")
    Y = tf.placeholder(tf.float32, shape=[None, n_classes], name="Y")
    keep_prob = tf.placeholder(tf.float32, name="keep_prob")

    # INPUT LAYER
    input_layer = tf.reshape(X, shape=[-1, img_rows, img_cols, img_depth])

    # CONVOLUTION LAYER
    conv_layer_1, conv_weights_1, conv_biases_1 = convolution_layer(
        input=input_layer,
        n_channels=1,
        filter_size=3,
        n_filters=32,
        name="conv_1"
    )

    # RELU ACTIVATION
    relu_layer_2 = relu_layer(conv_layer_1)

    # BATCH NORMALIZATION LAYER
    batch_norm_layer_3 = batch_normalization_layer(relu_layer_2)

    # MAX POOLING LAYER (2X2)
    max_pool_layer_4 = pooling_layer(batch_norm_layer_3, kernel_size=2, stride_size=2)

    # FLATTEN LAYER
    flatten_layer_5 = flatten_layer(max_pool_layer_4)

    # FULLY CONNECTED LAYER
    full_connect_layer_6, fc_weights_6, fc_biases_6 = fully_connected_layer(
        input=flatten_layer_5,
        n_inputs=6272, # calculated dimension of flattened_layer_5 = 14x14x32
        n_outputs=784,
        name="fc_layer_1"
    )

    # APPLY DROPOUT
    drop_out = tf.nn.dropout(full_connect_layer_6, keep_prob=keep_prob)

    # RELU ACTIVATION
    relu_layer_7 = relu_layer(drop_out)

    # FULLY CONNECTED LAYER
    full_connect_layer_8, fc_weights_8, fc_biases_8 = fully_connected_layer(

```

```

    input=relu_layer_7,
    n_inputs=784,
    n_outputs=10,
    name="fc_layer_2"
)

# SOFTMAX OUTPUT
Y_pred = tf.nn.softmax(full_connect_layer_8)

# CROSS ENTROPY LOSS
cross_entropy = tf.nn.softmax_cross_entropy_with_logits_v2(
    labels=Y,
    logits=full_connect_layer_8
)
# REGULARIZATION
regularizer = tf.nn.l2_loss(conv_weights_1) \
    + tf.nn.l2_loss(fc_weights_6) \
    + tf.nn.l2_loss(fc_weights_8)
# TOTAL LOSS
loss = tf.reduce_mean(cross_entropy) + weight_decay * regularizer

# CLASSIFICATION ACCURACY
Y_class = tf.argmax(Y, axis=1)
Y_pred_class = tf.argmax(Y_pred, axis=1)
correct_predictions = tf.equal(Y_class, Y_pred_class)
accuracy = tf.reduce_mean(tf.cast(correct_predictions, tf.float32))

# ADAM OPTIMIZER
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate).minimize(loss)

return optimizer, X, Y, keep_prob, loss, accuracy

```

```

def train_CNN(X_data, Y_data, batch_size, n_epochs, learning_rate, weight_decay, p):
    n_datapoints = X_data.shape[0]
    iter_per_epoch = int(np.ceil(n_datapoints/batch_size))
    iterations = iter_per_epoch * n_epochs

    # Create Loss/Accuracy dictionaries to store performance data
    loss_curves = {'train': [], 'valid': [], 'test': []}
    accuracy_curves = {'train': [], 'valid': [], 'test': []}

    # SET UP COMPUTATIONAL GRAPH
    optimizer, X, Y, keep_prob, loss, accuracy = build_CNN(learning_rate, weight_decay)
    global_init = tf.global_variables_initializer()

    with tf.Session() as sess:
        sess.run(global_init)
        for iter in range(iterations):
            # NEW EPOCH
            if iter % iter_per_epoch == 0:
                print("EPOCH {}".format(int((iter+1)/iter_per_epoch)))

            # CALCULATE TRAINING LOSS & ACCURACY
            feed_dict_train = {X: X_data, Y: Y_data, keep_prob: 1.0}
            _loss, _acc = sess.run([loss, accuracy], feed_dict=feed_dict_train)
            loss_curves['train'].append(_loss)
            accuracy_curves['train'].append(_acc)

            # CALCULATE VALIDATION LOSS & ACCURACY
            feed_dict_valid = {X: X_valid, Y: Y_valid, keep_prob: 1.0}
            _loss, _acc = sess.run([loss, accuracy], feed_dict=feed_dict_valid)

```

```

loss_curves['valid'].append(_loss)
accuracy_curves['valid'].append(_acc)

# CALCULATE TEST LOSS & ACCURACY
feed_dict_test = {X: X_test, Y: Y_test, keep_prob: 1.0}
_loss, _acc = sess.run([loss, accuracy], feed_dict=feed_dict_test)
loss_curves['test'].append(_loss)
accuracy_curves['test'].append(_acc)

# SHUFFLE DATA ON NEW EPOCH
X_data, Y_data = shuffle(X_data, Y_data)

# SELECT MINI-BATCH
X_batch, Y_batch = X_data[:batch_size], Y_data[:batch_size]

# GRADIENT DESCENT STEP on mini-batch
feed_dict_batch = {X: X_batch, Y: Y_batch, keep_prob: p}
sess.run([optimizer], feed_dict=feed_dict_batch)

# SHIFT DATA BY BATCH SIZE so next sample is new
X_data = np.roll(X_data, batch_size, axis=0)
Y_data = np.roll(Y_data, batch_size, axis=0)

```

```

return loss_curves, accuracy_curves

```

```

def main():
    start_time = time.time()
    loss, accuracy = train_CNN(
        X_train,
        Y_train,
        batch_size=32,
        n_epochs=50,
        learning_rate=1e-4,
        weight_decay=0,
        p=1.0
    )
    end_time = time.time()
    print("--- %s seconds ---" % (time.time() - start_time))

    print("TRAINING -----")
    print("Loss: ", loss['train'][-1])
    print("Accuracy:", accuracy['train'][-1])

    print("VALIDATION -----")
    print("Loss: ", loss['valid'][-1])
    print("Accuracy:", accuracy['valid'][-1])

    print("TESTING -----")
    print("Loss: ", loss['test'][-1])
    print("Accuracy:", accuracy['test'][-1])

    plt.plot(loss['train'], color='blue', label='training data')
    plt.plot(loss['valid'], color='red', label='validation data')
    plt.plot(loss['test'], color='green', label='test data')
    plt.legend()
    plt.title('Loss Curves')
    plt.ylabel('CE Loss')
    plt.xlabel('Epoch')
    plt.show()

```

```
plt.plot(accuracy['train'], color='blue', label='training data')
plt.plot(accuracy['valid'], color='red', label='validation data')
plt.plot(accuracy['test'], color='green', label='test data')
plt.legend()
plt.title('Accuracy Curves')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.show()
```

```
if __name__ == '__main__':
    main()
```